

# CSC139 Operating System Principles

Fall 2020, Part 3-1

Instructor: Dr. Yuan Cheng


# Where we are in the course

- Discussed
  - Processes & Threads
  - CPU Scheduling
  - Synchronization & Deadlock
- Next up
  - Memory Management
- Yet to come
  - File Systems & I/O Storage
  - Security and Protection

# Memory Management


- Where is the executing process?
- How do we allow multiple processes to use main memory simultaneously?
- What is an address and how is one interpreted?

# Background



- Program executable starts out on disk
- The OS loads the program into memory
- CPU fetches instructions and data from memory while executing the program

# Terminology



- Segment: a chunk of memory assigned to a process
- Physical Address: a real address in memory
- Virtual Address: an address relative to the start of a process's address space

# Where do addresses come from?

- How do programs generate instruction and data addresses?
  - **Compile time**: The compiler generates the exact physical location in memory starting from some fixed starting position  $k$ . The OS does nothing.
  - **Load time**: The compiler generates an address, but at load time the OS determines the process's starting position. Once the process loads, it does not move in memory.
  - **Execution time**: The compiler generates an address, and the OS can place it anywhere it wants in memory.

# Logical vs. Physical Address Space


- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program

# Uni-programming

- OS gets a fixed part of memory (highest memory in DOS).
- One process executes at a time.
- Process is always loaded starting at address 0.
- Process executes in a contiguous section of memory.
- Compiler can generate physical addresses.
- Maximum address = Memory Size - OS Size
- OS is protected from process by checking addresses used by process.



# Uni-programming



=> Simple, but does not allow for overlap of I/O and computation

# Multiple Programs Share Memory

- Transparency:
  - We want multiple processes to coexist in memory.
  - No process should be aware that memory is shared.
  - Processes should not care what physical portion of memory they are assigned to.
- Safety:
  - Processes must not be able to corrupt each other.
  - Processes must not be able to corrupt the OS.
- Efficiency:
  - Performance of CPU and memory should not be degraded badly due to sharing.

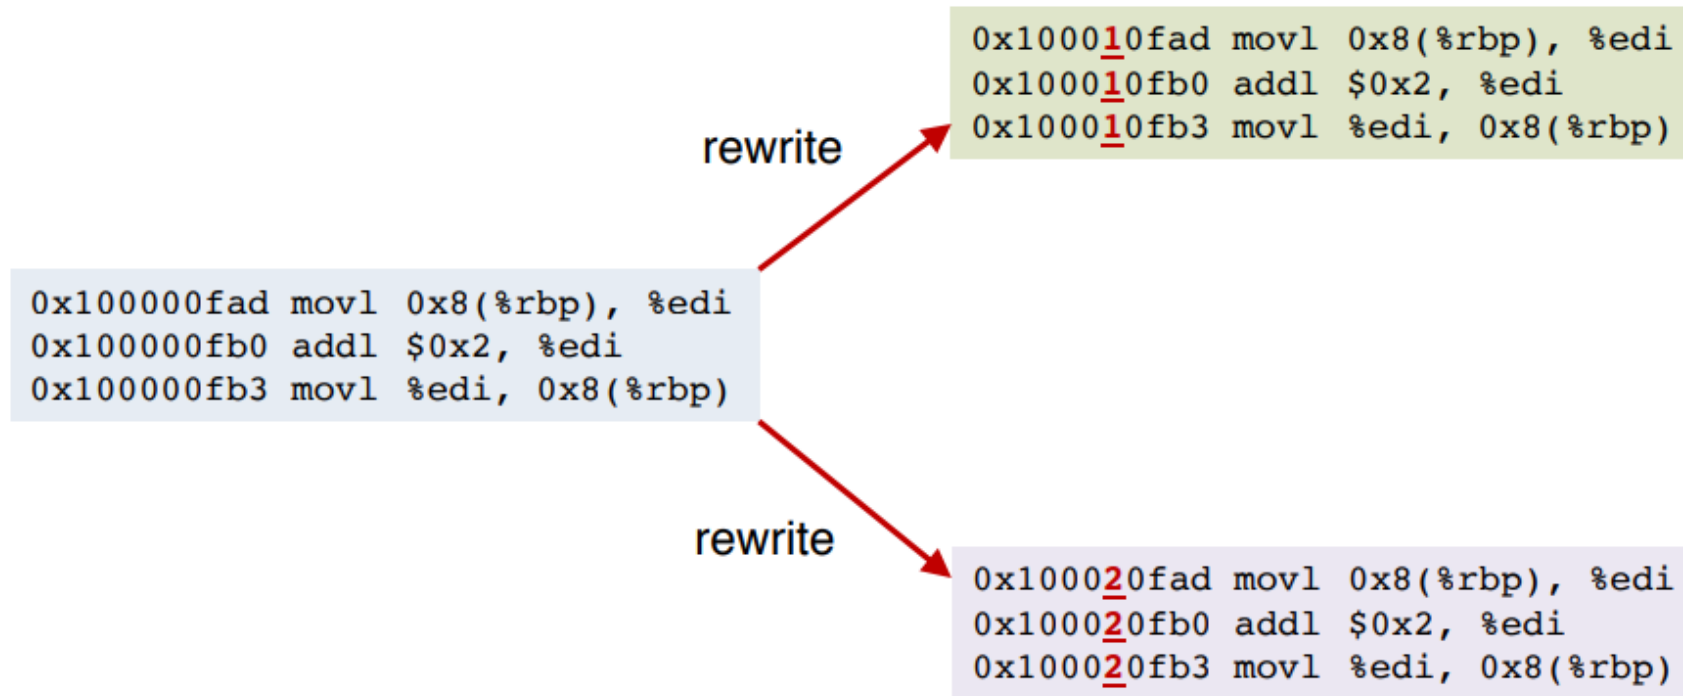
# How to Run Multiple Programs?

- Approaches:
  - Static relocation
  - Dynamic relocation
  - Segmentation

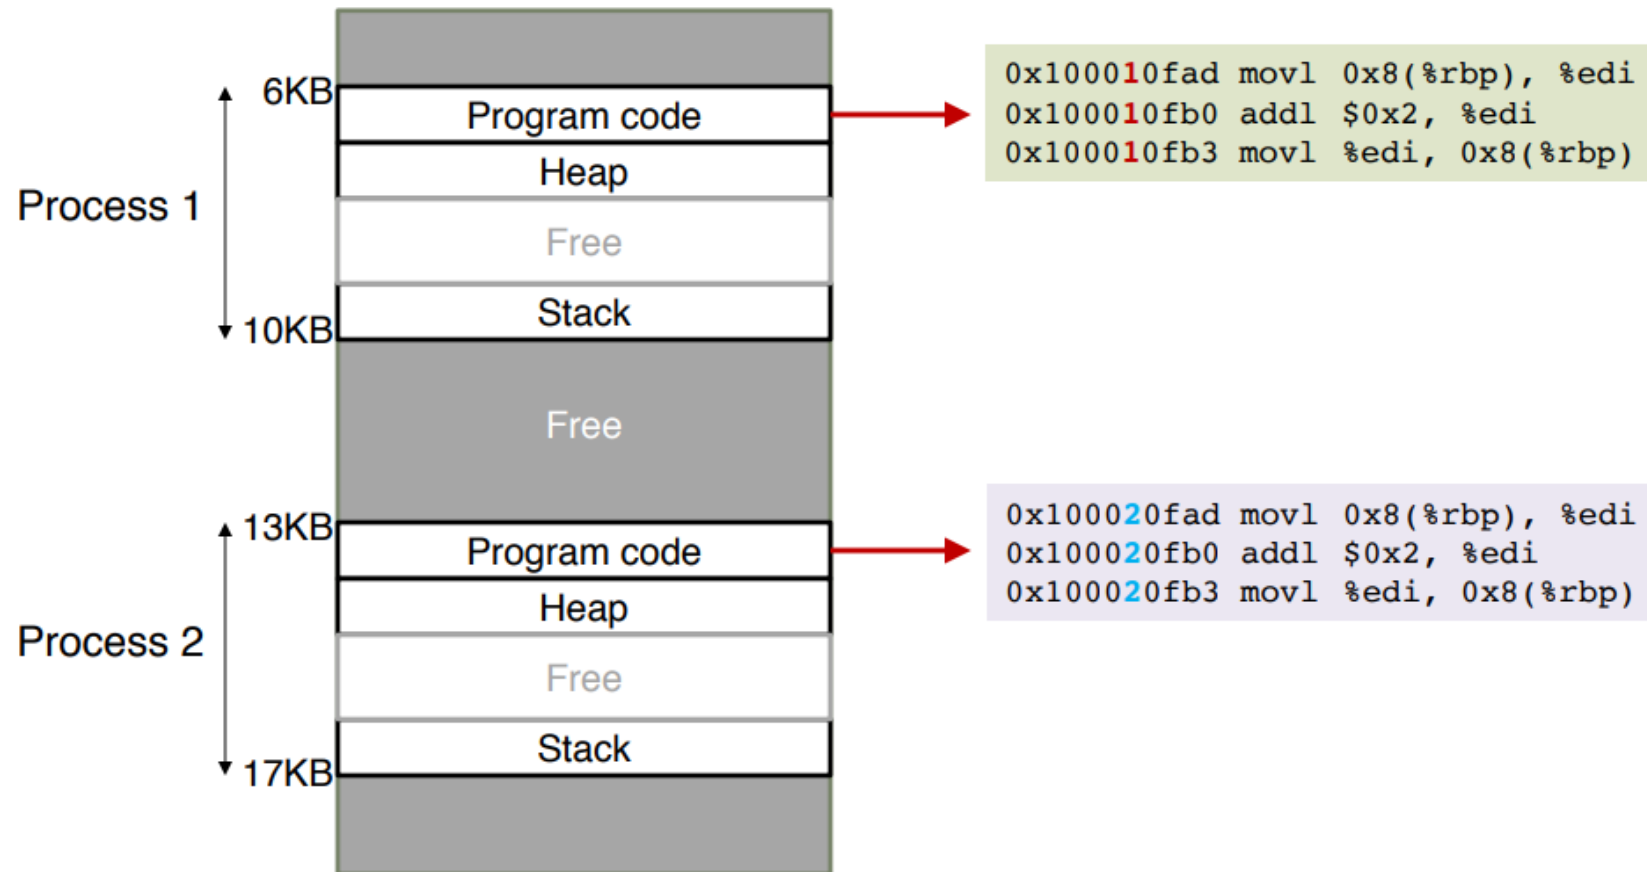
# Static Relocation

- At load time, the OS adjusts the addresses in a process to reflect its position in memory.
  - Each rewrite uses **different** addresses and pointers
  - Change jumps, loads, etc.
- Once a process is assigned a place in memory and starts executing it, the OS cannot move it. (Why?)

# Rewrite for Each New Process



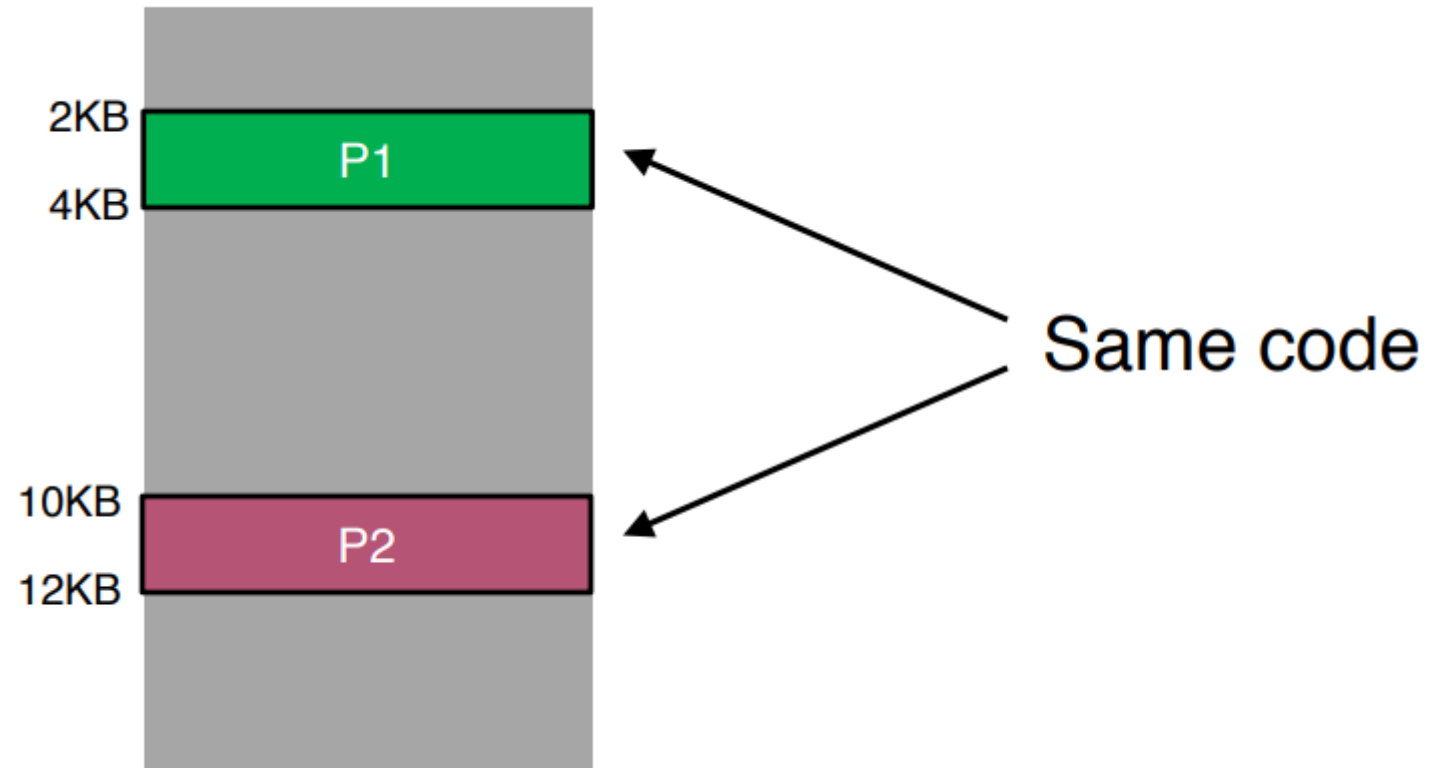
# Rewrite for Each New Process



# Dynamic Relocation

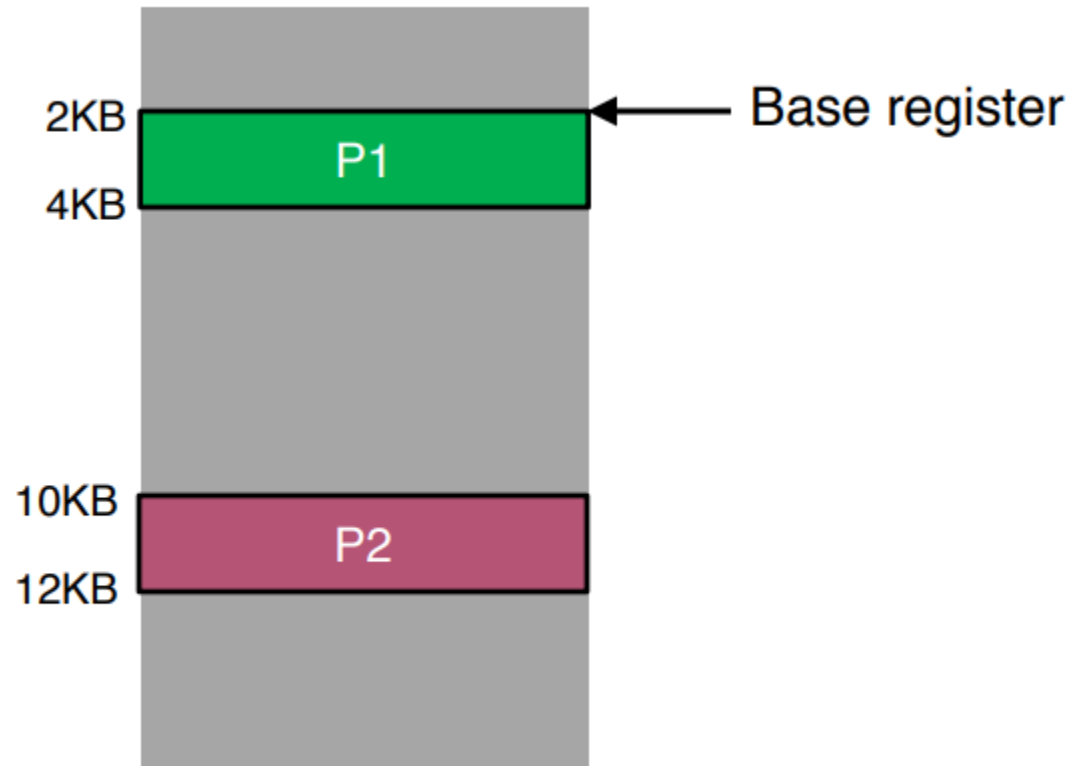
- Hardware adds relocation register (base) to virtual address to get a physical address;
  - Each process has a **different** value in the base register when running
- Hardware compares address with limit register (address must be less than limit).
- If test fails, the processor takes an address trap and ignores the physical address.

# Base Relocation



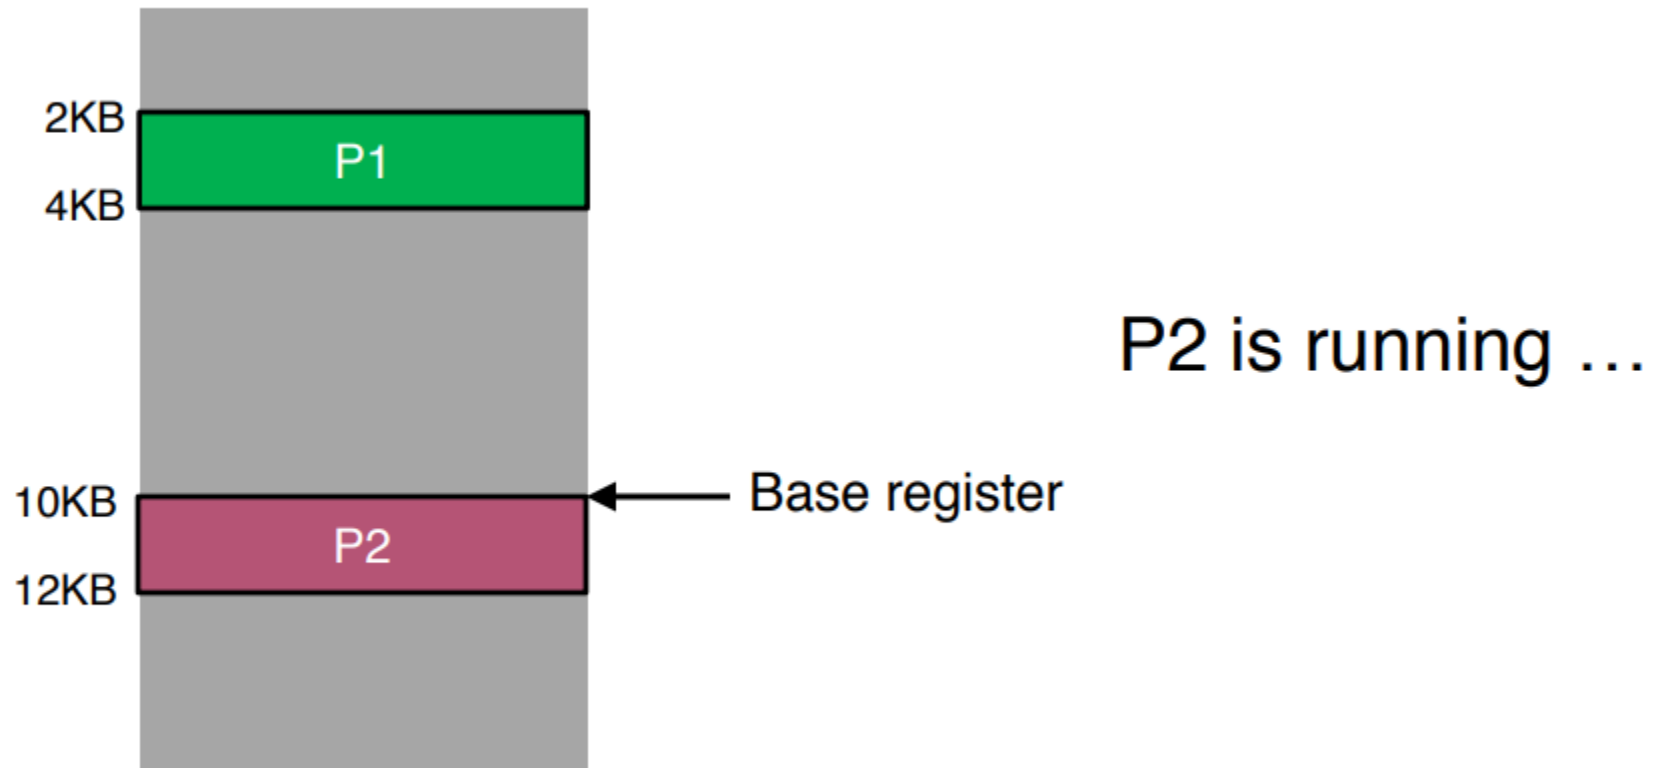


# Base Relocation




P1 is running ...

# Base Relocation




# Base Relocation




Virtual	Physical
P1: load 100, R1	load 2148, R1

# Base Relocation




Virtual	Physical
P1: load 100, R1	load 2148, R1
P2: load 100, R1	load 10340, R1

# Base Relocation




Virtual	Physical
P1: load 100, R1	load 2148, R1
P2: load 100, R1	load 10340, R1
P2: load 2000, R1	load 12240, R1

# Base Relocation



Virtual	Physical
P1: load 100, R1	load 2148, R1
P2: load 100, R1	load 10340, R1
P2: load 2000, R1	load 12240, R1
P1: load 2000, R1	load 4048, R1

# Base Relocation




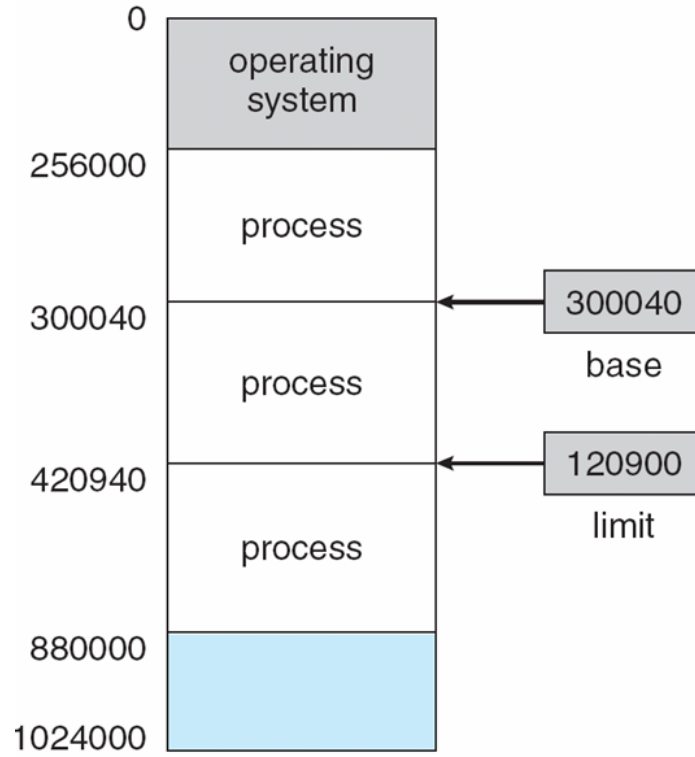
Virtual	Physical
P1: load 100, R1	load 2148, R1
P2: load 100, R1	load 10340, R1
P2: load 2000, R1	load 12240, R1
P1: load 2000, R1	load 4048, R1
<b>P1: store 9241, R1</b>	<b>store 11289, R1</b>

**Can P1 hurt P2?**

**Can P2 hurt P1?**

# Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user





# Dynamic Relocation


- Advantages:
  - OS can easily move a process during execution.
  - OS can allow a process to grow over time.
  - Simple, fast hardware: two special registers, an add, and a compare.
- Disadvantages:
  - Slows down hardware due to the add on every memory reference.
  - Can't share memory (such as program text) between processes.
  - Process is still limited to physical memory size.
  - Degree of multiprogramming is very limited since all memory of all active processes must fit in memory.
  - Complicates memory management.

# Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
  - Resident operating system, usually held in low memory with interrupt vector
  - User processes then held in high memory
  - Each process contained in single contiguous section of memory

# Memory Allocation

- As processes enter the system, grow, and terminate, the OS must keep track of which memory is available and utilized



- Holes: pieces of free memory (shaded area in the above figure)
- Given a new process, the OS must decide which hole to use for the process

# Memory Allocation Policies

- **First-Fit**: allocate the first one in the list in which the process fits. The search can start with the first hole, or where the previous first-fit search ended.
- **Best-Fit**: Allocate the smallest hole that is big enough to hold the process. The OS must search the entire list or store the list sorted by size hole list.
- **Worst-Fit**: Allocate the largest hole to the process. Again the OS must search the entire list or keep the list sorted.
- Simulations show first-fit and best-fit usually yield better storage utilization than worst-fit; first-fit is generally faster than best-fit

# External Fragmentation

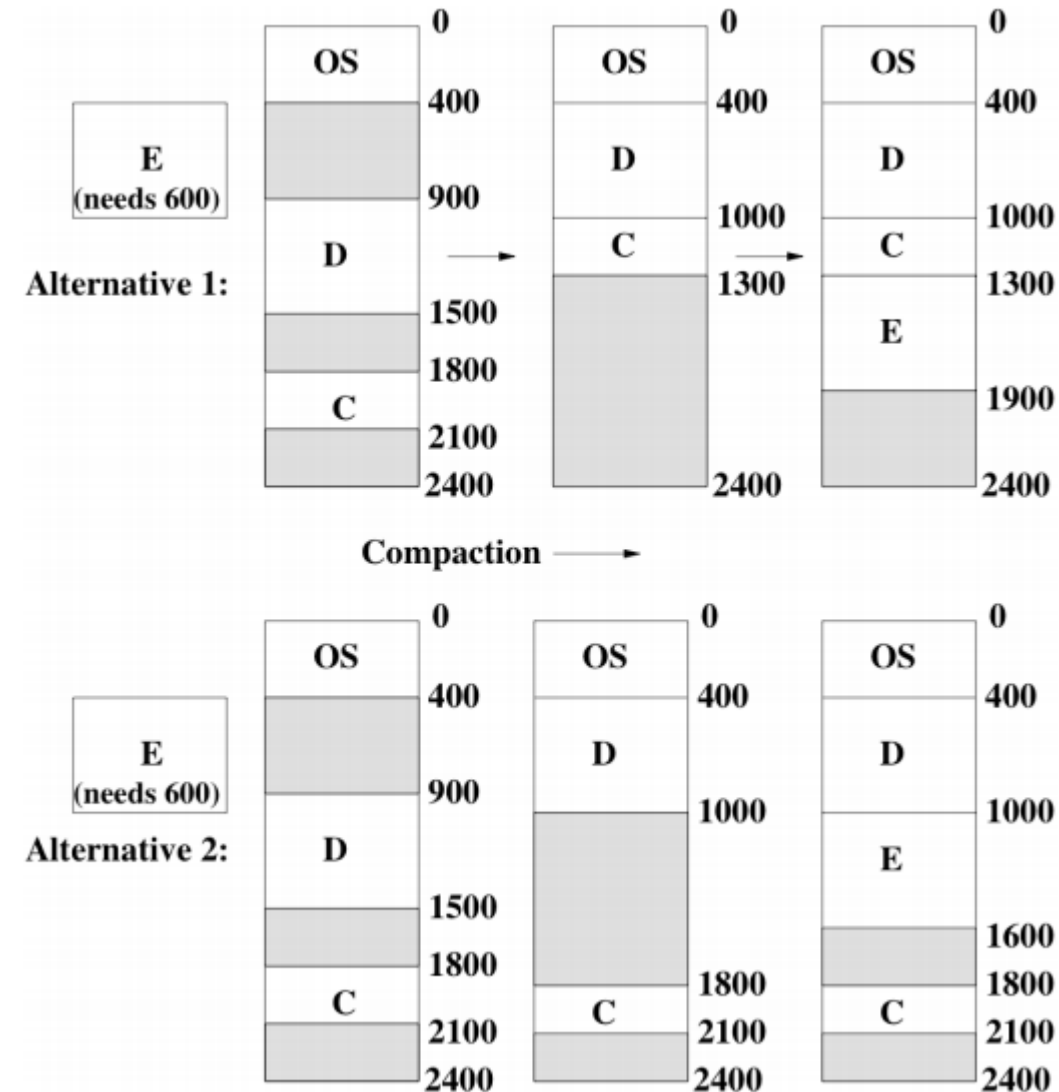
- Frequent loading and unloading programs causes free space to be broken into little pieces
- External fragmentation exists when there is enough memory to fit a process in memory, but the space is not contiguous
- 50-percent rule: Simulations show that for every  $2N$  allocated blocks,  $N$  blocks are lost due to fragmentation (i.e.,  $1/3$  of memory space is wasted)
- We want an allocation policy that minimizes wasted space.

# Internal Fragmentation

- Consider a process of size 8846 bytes and a block of size 8848 bytes  
⇒ it is more efficient to allocate the process the entire 8848 block than it is to keep track of 2 free bytes
- Internal fragmentation exists when memory internal to a partition that is wasted

# Compaction

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time
- How much memory is moved?
- How big a block is created?
- Any other choices?




# Swapping

- Roll out a process to disk, releasing all the memory it holds.
- When process becomes active again, the OS must reload it in memory.
  - With static relocation, the process must be put in the same position.
  - With dynamic relocation, the OS finds a new position in memory for the process and updates the relocation and limit registers.
- If swapping is part of the system, compaction is easy to add.
- How could or should swapping interact with CPU scheduling?



# Systematic View of Swapping



# Problems

- Fragmentation
  - Frequent compaction needed
- Contiguous allocation
  - Difficult to grow or shrink process memory
- Requirement that process resides entirely in memory
  - Swapping helps but not perfect


# Session Plan

- Segmentation
- Paging


# Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
  - A segment is a logical unit such as:
    - main program
    - procedure
    - function
    - method
    - object
    - local variables, global variables
    - common block
    - stack
    - symbol table
    - arrays


# User's View of a Program



# Logical View of Segmentation



user space



physical memory space

# Segmentation Architecture

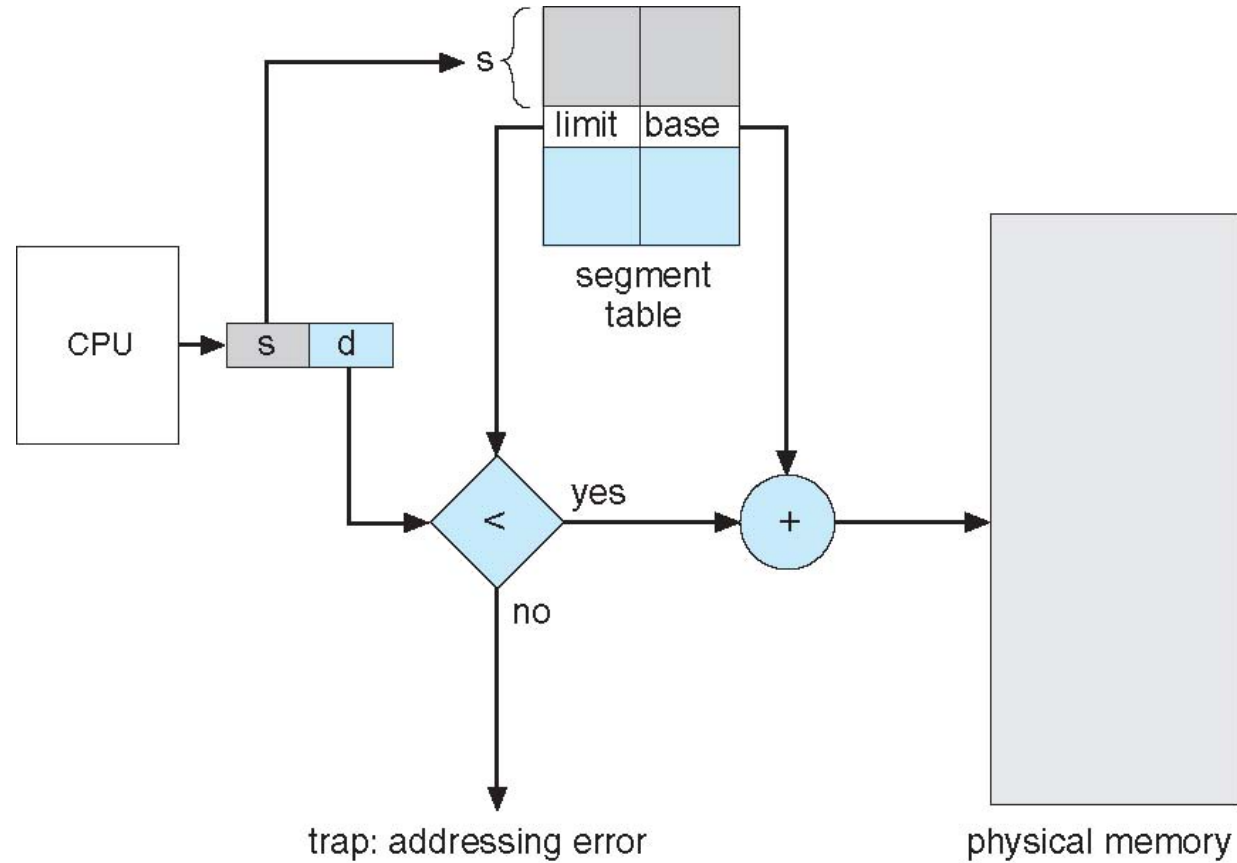
- Logical address consists of a two tuple:  
    <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
  - **base** – contains the starting physical address where the segments reside in memory
  - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;  
    segment number **s** is legal if **s** < **STLR**

# Segmentation Architecture (cont.)

- Protection
  - With each entry in segment table associate:
    - validation bit = 0  $\Rightarrow$  illegal segment
    - read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram



# Segmentation Hardware



# Paging: Motivation & Features

- 90/10 rule: Processes spend 90% of their time accessing 10% of their space in memory

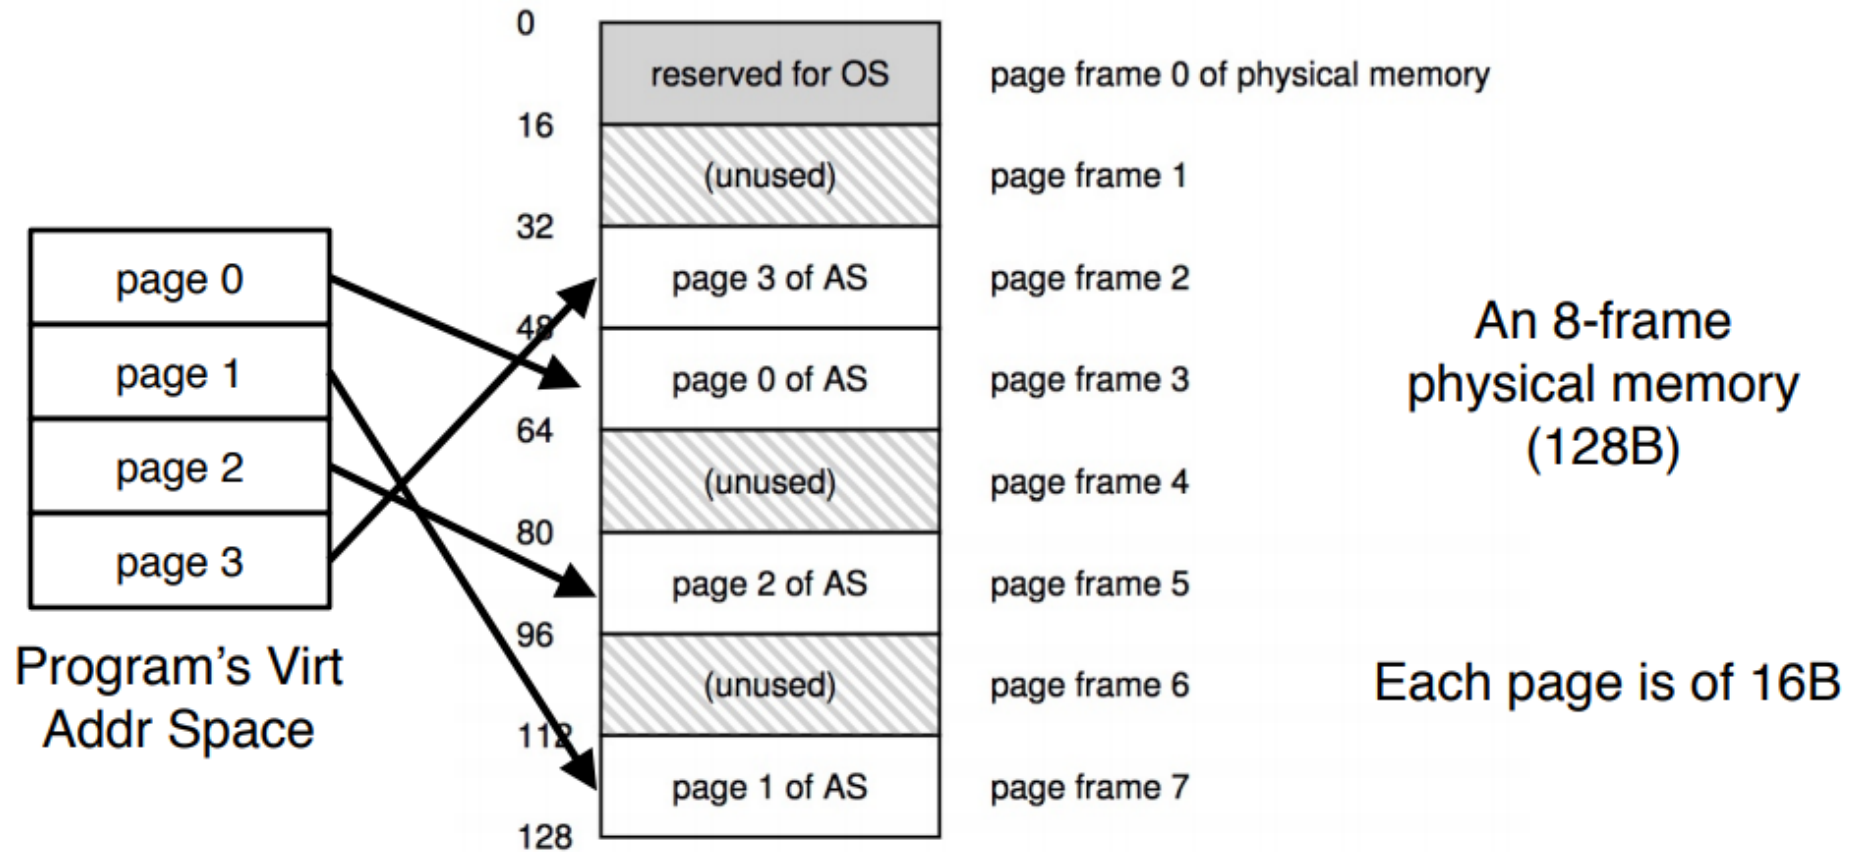
=> Keep only those parts of a process in memory that are actually being used

- Pages greatly simplify the hole fitting problem
- The logical memory of the process is contiguous, but pages need not be allocated contiguously in memory
- By dividing memory into fixed size pages, we can eliminate external fragmentation
- Paging does not eliminate internal fragmentation

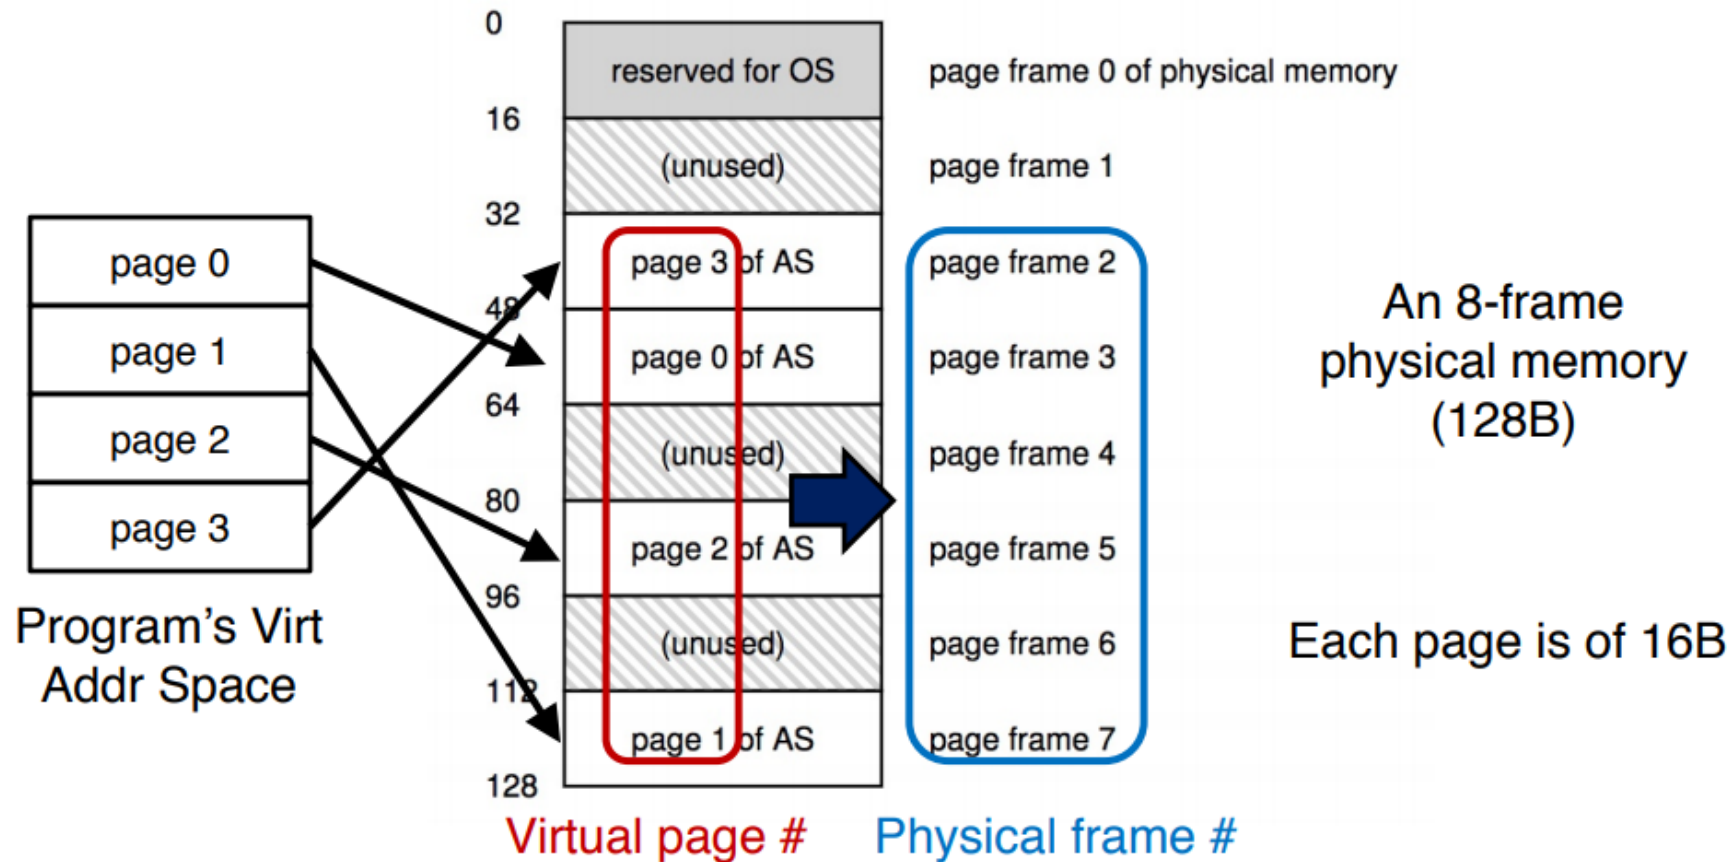
# Paging

- A memory management scheme that allows the physical address space of a process to be **non-contiguous**
- Divide **physical memory** into fixed-sized blocks called **frames**
- Divide **logical memory** into blocks of same size called **pages**
- Flexible mapping: Any page can go to any free frame
- Scalability: To run a program of size  $n$  pages, need to find  $n$  free frames and load program

# A Simple Example

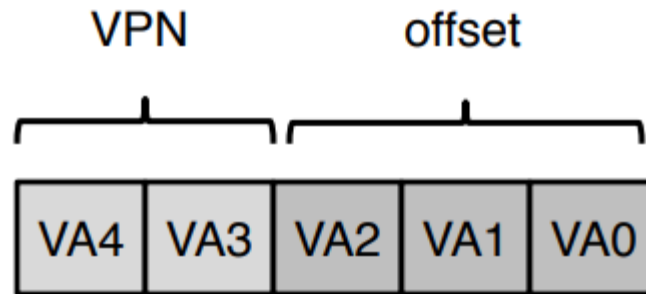


# A Simple Example (cont.)



# Addressing Basics

- For segmentation
  - High bits => segment #
  - Low bits => offset
- For paging
  - High bits => page #
  - Low bits => offset



Q: How many offset bits do we need?

A:  $\log(\text{page\_size})$

# Address Examples

Page size	Offset
16 Bytes	
2 KB	
4 MB	
256 Bytes	
16 KB	

# Address Examples

Page size	Offset
16 Bytes	4
2 KB	11
4 MB	22
256 Bytes	8
16 KB	14



# Address Examples

Page size	Offset	Virtual Address Bits	VPN
16 Bytes	4	10	
2 KB	11	20	
4 MB	22	32	
256 Bytes	8	16	
16 KB	14	64	

# Address Examples

Page size	Offset	Virtual Address Bits	VPN
16 Bytes	4	10	6
2 KB	11	20	9
4 MB	22	32	10
256 Bytes	8	16	8
16 KB	14	64	50

# Address Examples

Page size	Offset	Virtual Address Bits	VPN	# of Virtual Pages
16 Bytes	4	10	6	
2 KB	11	20	9	
4 MB	22	32	10	
256 Bytes	8	16	8	
16 KB	14	64	50	

# Address Examples


Page size	Offset	Virtual Address Bits	VPN	# of Virtual Pages
16 Bytes	4	10	6	64
2 KB	11	20	9	512
4 MB	22	32	10	1K
256 Bytes	8	16	8	256
16 KB	14	64	50	$2^{50}$

# Page Table

- A **per-process** data structure used to keep track of virtual page to physical frame mapping
- Major role: store **address translation**

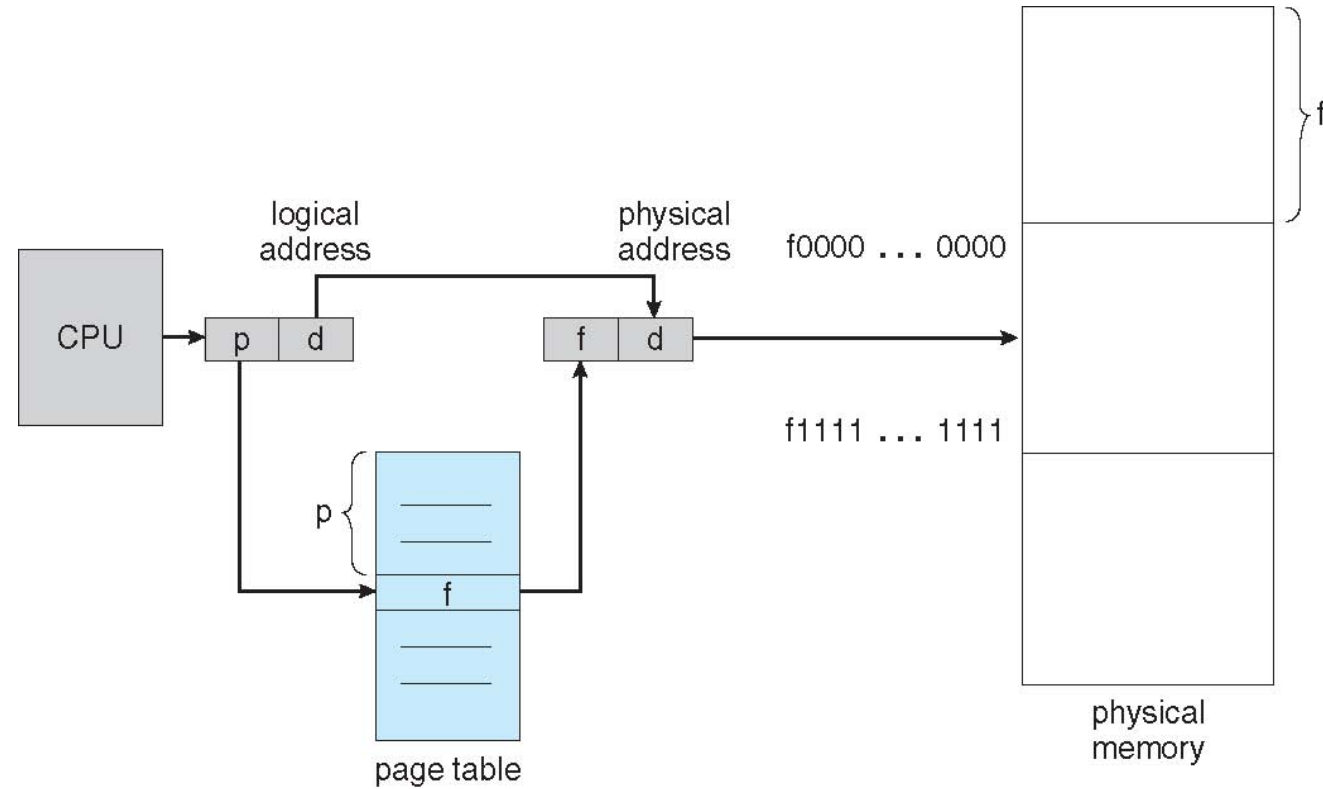
# Address Translation Scheme

- m-bit virtual address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a **page table** which contains base address of each page in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit



- For given logical address space  $2^m$  and page size  $2^n$

# Paging Hardware

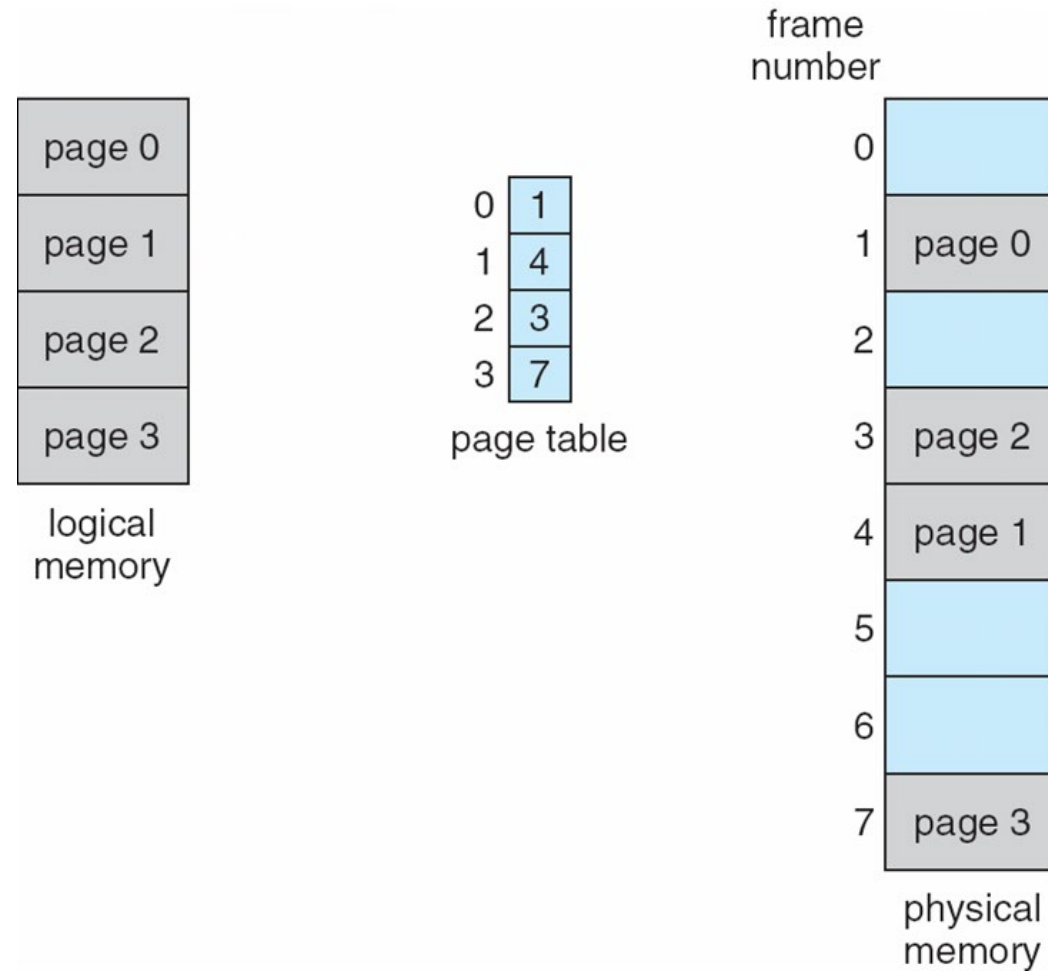


# Paging Hardware


- Paging is a form of dynamic relocation, where each virtual address is bound by the paging hardware to a physical address.
- Think of the page table as a set of relocation registers, one for each frame.
- Mapping is invisible to the process; the OS maintains the mapping and the hardware does the translation.
- Protection is provided with the same mechanisms as used in dynamic relocation.



# Paging Model of Logical and Physical Memory



# Paging Example




$n=2$  and  $m=4$  32-byte memory and 4-byte pages

# Paging (cont.)


- Calculating internal fragmentation
  - Page size = 2,048 bytes
  - Process size = 72,766 bytes
  - 35 pages + 1,086 bytes
  - Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
  - Worst case fragmentation = 1 frame – 1 byte
  - On average fragmentation =  $1 / 2$  frame size
  - So small frame sizes desirable?
  - But each page table entry takes memory to track
  - Page sizes growing over time
    - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

# Free Frames



(a)

Before allocation



(b)

After allocation


# Implementation of Page Table

- Page table is kept in main memory
- Each **page table entry (PTE)** holds  
    <physical translation + other info>
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table

# Page Table Entry (PTE)

- The simplest form of a page table is a **linear page table**
  - Array data structure
  - OS indexes the array by virtual page number (VPN)
  - To find the desired physical frame number (PFN)

## An 32-bit x86 page table entry (PTE)



# Paging Problems

- Page tables are too slow
- Page tables are too big

# Translation Lookaside Buffer (TLB)



# Performance Problems of Paging

- A basic memory access protocol
  1. Fetch the translation from in-memory page table
  2. Explicit load/store access on a memory address
- In this scheme every data/instruction access requires **two** memory accesses
  - One for the page table
  - And one for the data/instruction
- Too much performance overhead!

# Speeding Up Translation

- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
- A TLB is part of the memory-management unit (MMU)
- A TLB is a hardware cache
- Algorithm sketch
  - For each virtual memory reference, hardware first checks the TLB to see if the desired translation is held therein

# TLB Basic Algorithm

1. Extract VPN from VA
2. Check if TLB holds the translation
3. If it is a **TLB hit** – extract PFN from the TLB entry, concatenate it onto the offset to form the PA
4. If it is a **TLB miss** – access **page table** to get the translation, update the TLB entry with the translation

# How Many TLB Lookups?

- Assume 4KB pages

```
int sum = 0;
for (i=0; i<1024; i++) {
    sum += a[i];
}
```

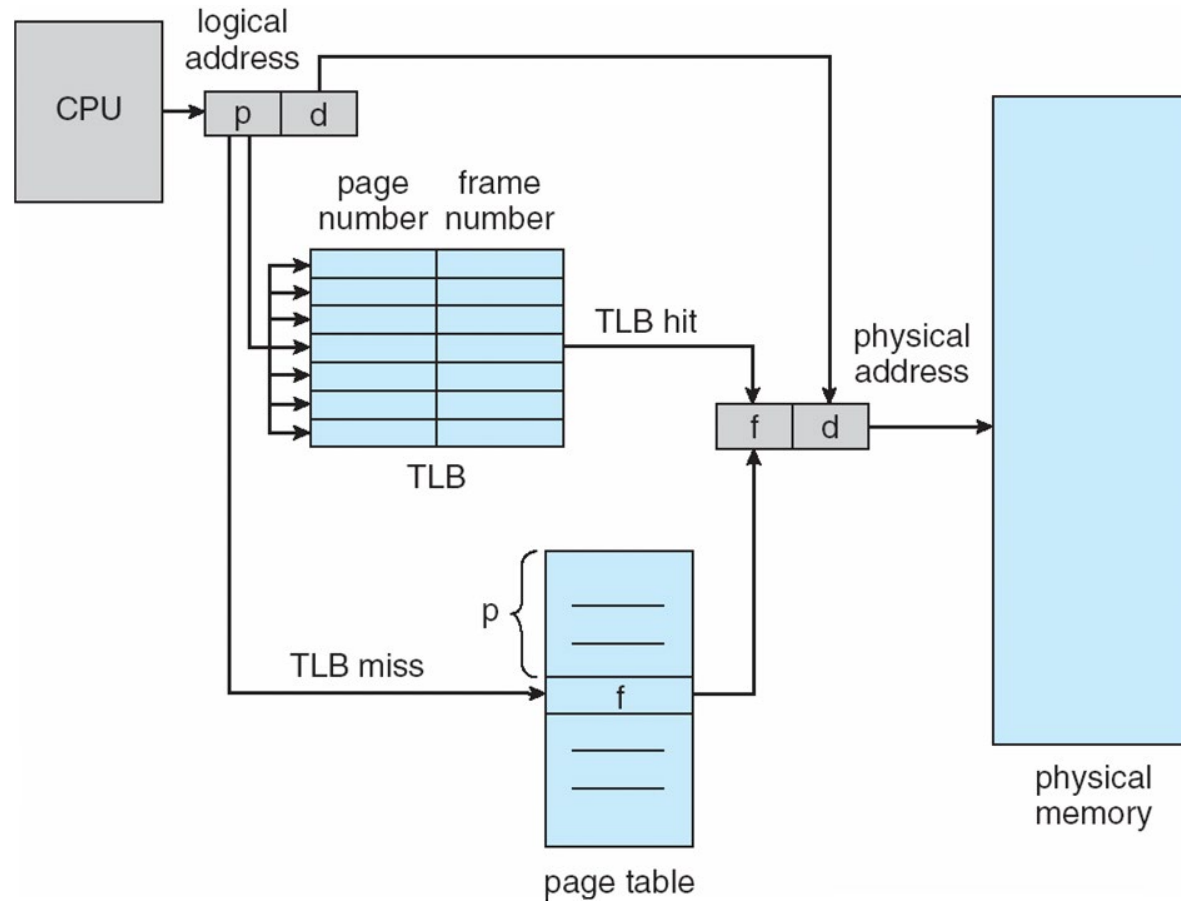
# How Many TLB Lookups?

- Assume 4KB pages

```
int sum = 0;
for (i=0; i<1024; i++) {
    sum += a[i];
}
```

- Array a has 1024 items, each item is 4 bytes:  $\text{size}(a) = 4096$  bytes
- Number of TLB miss =  $4096/4096 = 1$
- TLB miss rate =  $1/1024 = 0.09\%$

# Paging Hardware with TLB



# Effective Access Time

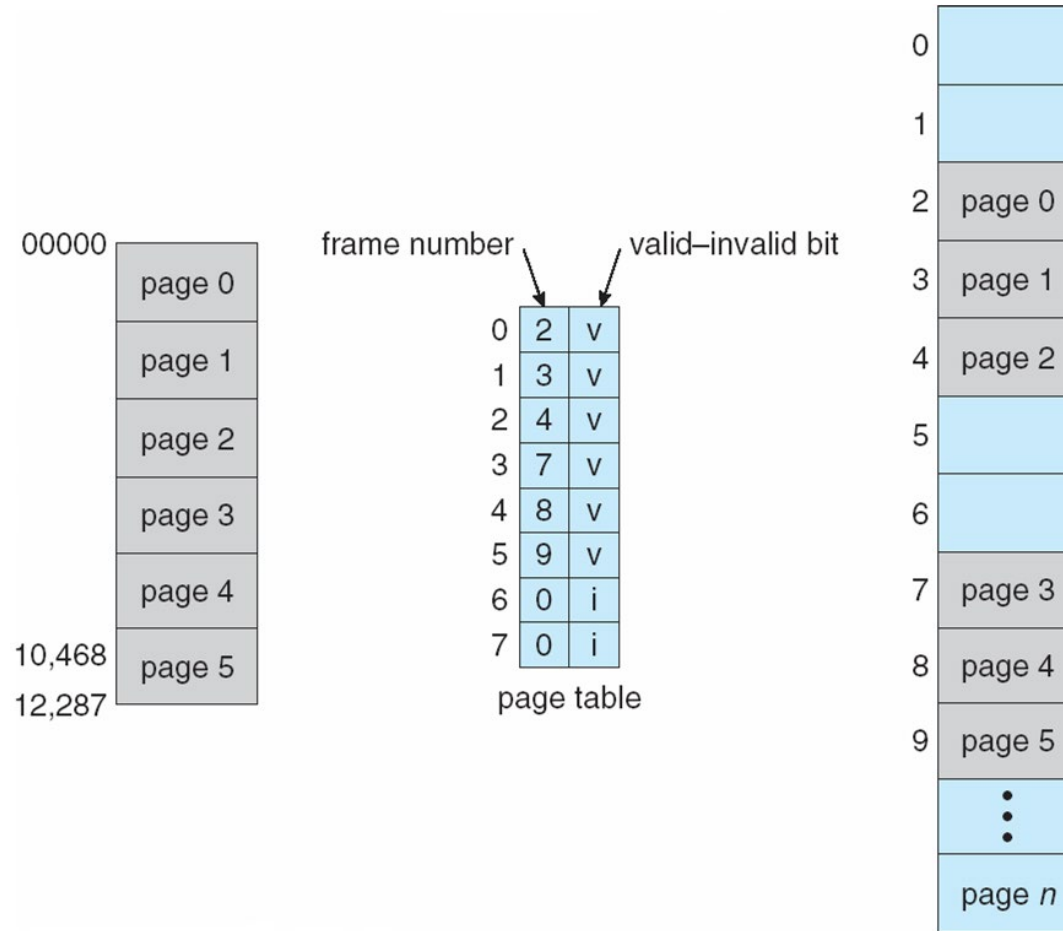
- Associative Lookup =  $\varepsilon$  time unit
  - Can be  $< 10\%$  of memory access time
- Hit ratio =  $\alpha$ 
  - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Consider  $\alpha = 80\%$ ,  $\varepsilon = 100\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio  $\rightarrow \alpha = 99\%$ ,  $\varepsilon = 100\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

# Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel



# Valid(v) or Invalid(i) Bit in a Page Table




# Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers
  - Page size of 4 KB ( $2^{12}$ )
  - Page table would have 1 million entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
    - That amount of memory used to cost a lot
    - Don't want to allocate that contiguously in main memory
- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

# Hierarchical Page Tables


- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table

# Two-Level Page-Table Scheme



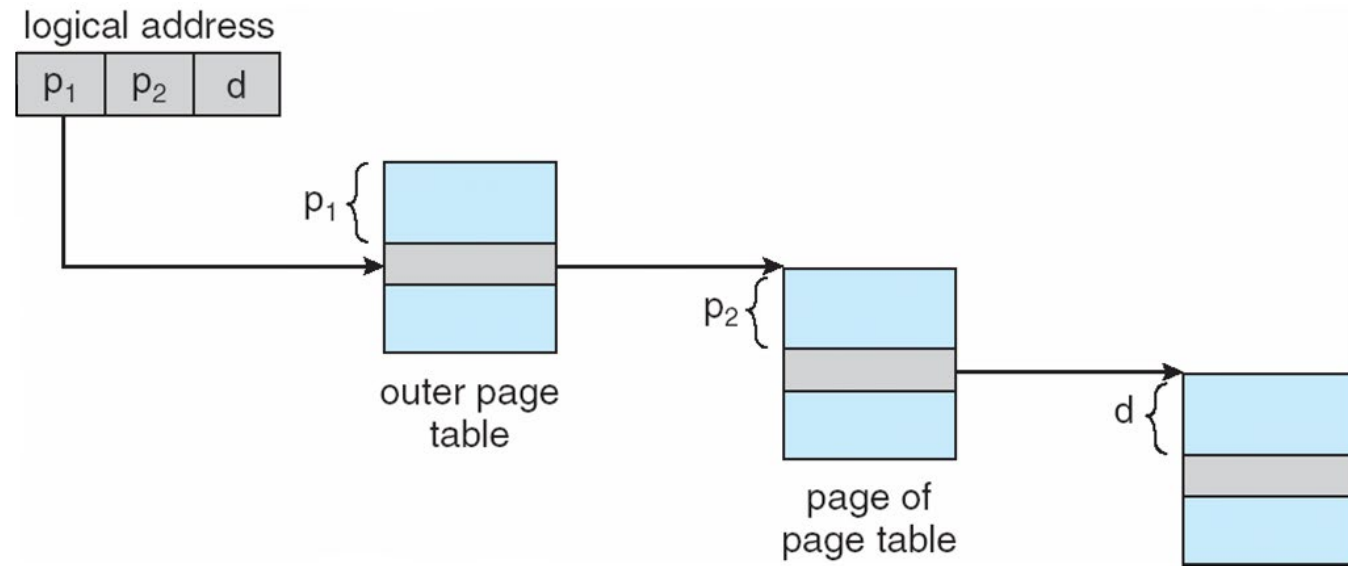
# Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:




- where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

# Address-Translation Scheme





# 64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like



- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes
- One solution is to add a  $2^{\text{nd}}$  outer page table
- But in the following example the  $2^{\text{nd}}$  outer page table is still  $2^{34}$  bytes in size
  - And possibly 4 memory access to get to one physical memory location

# Three-Level Paging Scheme






# Hashed Page Tables

- Common in address spaces  $> 32$  bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
  - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)

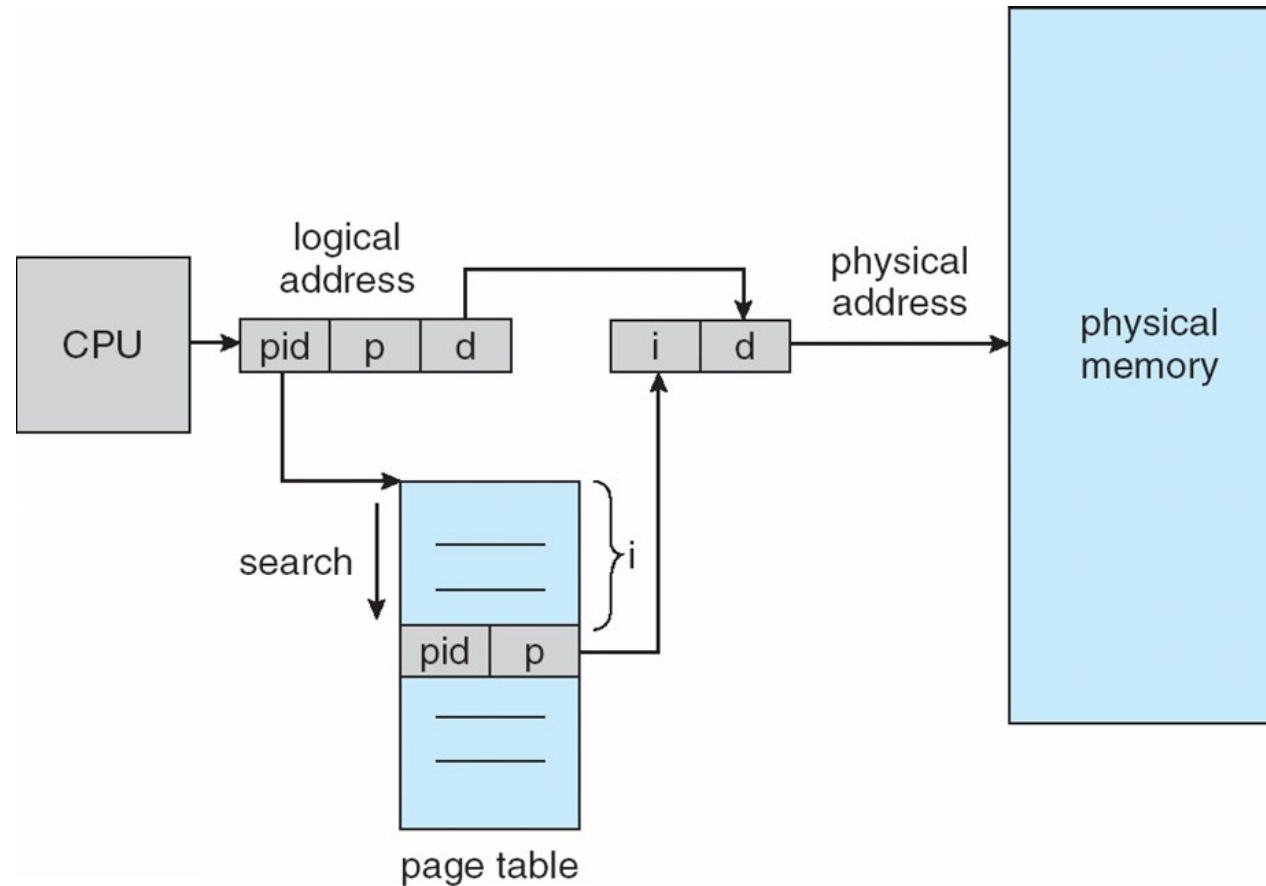
# Hashed Page Table



# Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
  - TLB can accelerate access
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address

# Inverted Page Table Architecture



# Exit Slips

- Take 1-2 minutes to reflect on this lecture
- On a sheet of paper write:
  - One thing you learned in this lecture
  - One thing you didn't understand

# Next class

- We will continue to discuss:
  - Virtual Memory
- Reading assignment:
  - SGG: Ch. 10

# Acknowledgment

- The slides are partially based on the ones from
  - The book site of *Operating System Concepts (Tenth Edition)*: <http://os-book.com/>